

# Kanister: Application-Level Data Operations on Kubernetes

Michael Cade  
Global Field CTO  
Veeam Software



# Michael Cade

Michael Cade is a seasoned technology executive with over 20 years of experience in the data protection space. Currently serving as the Global Field CTO for Veeam Software.



90DaysOfDevOps

 @MichaelCade1

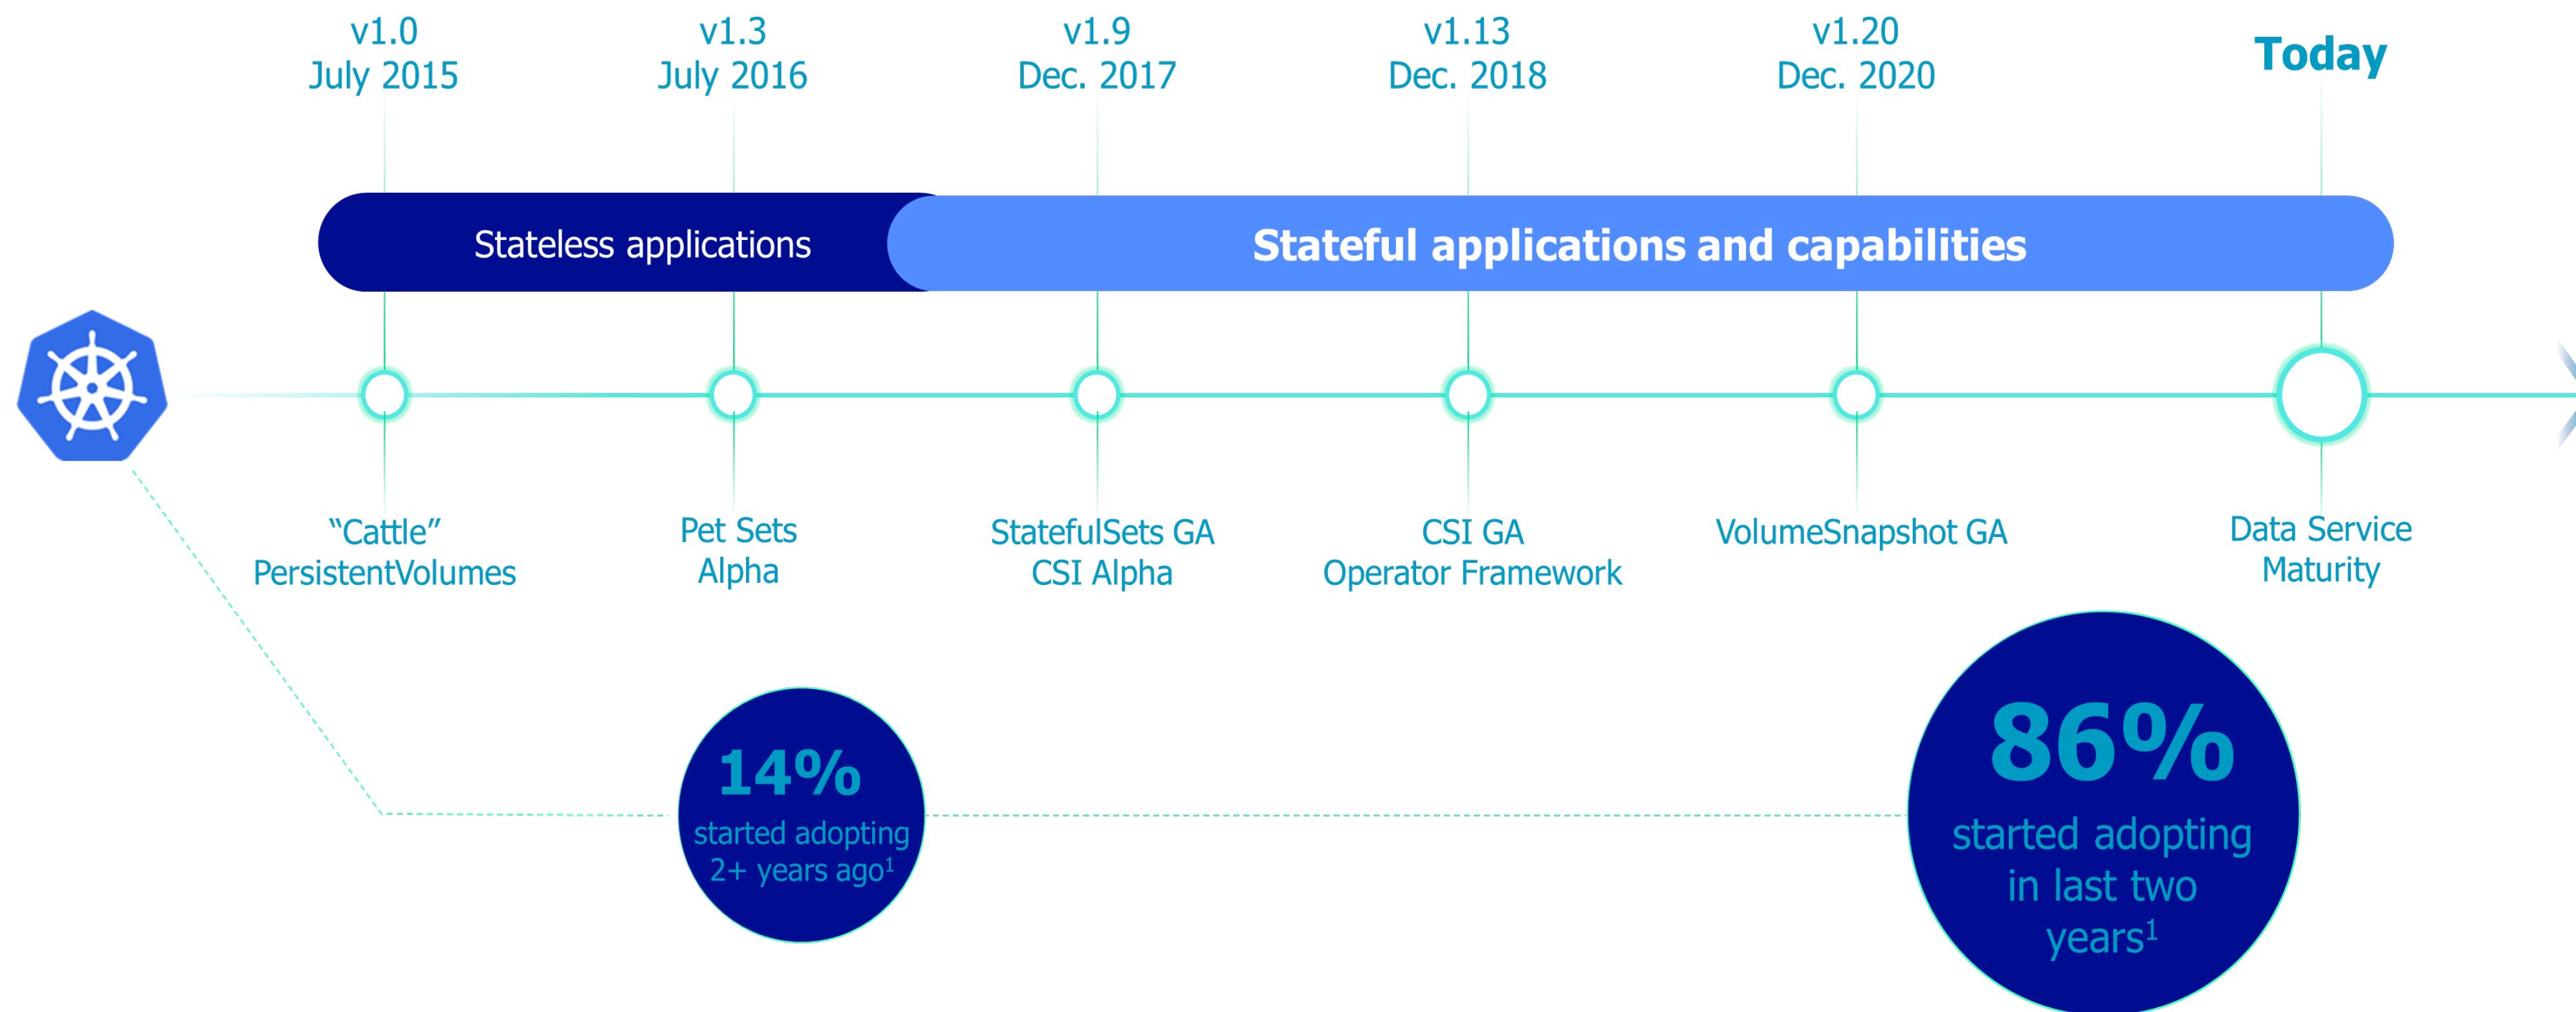
 YouTube

 LinkedIn

Michael Cade  
Field CTO, Cloud Native  
Veeam Software

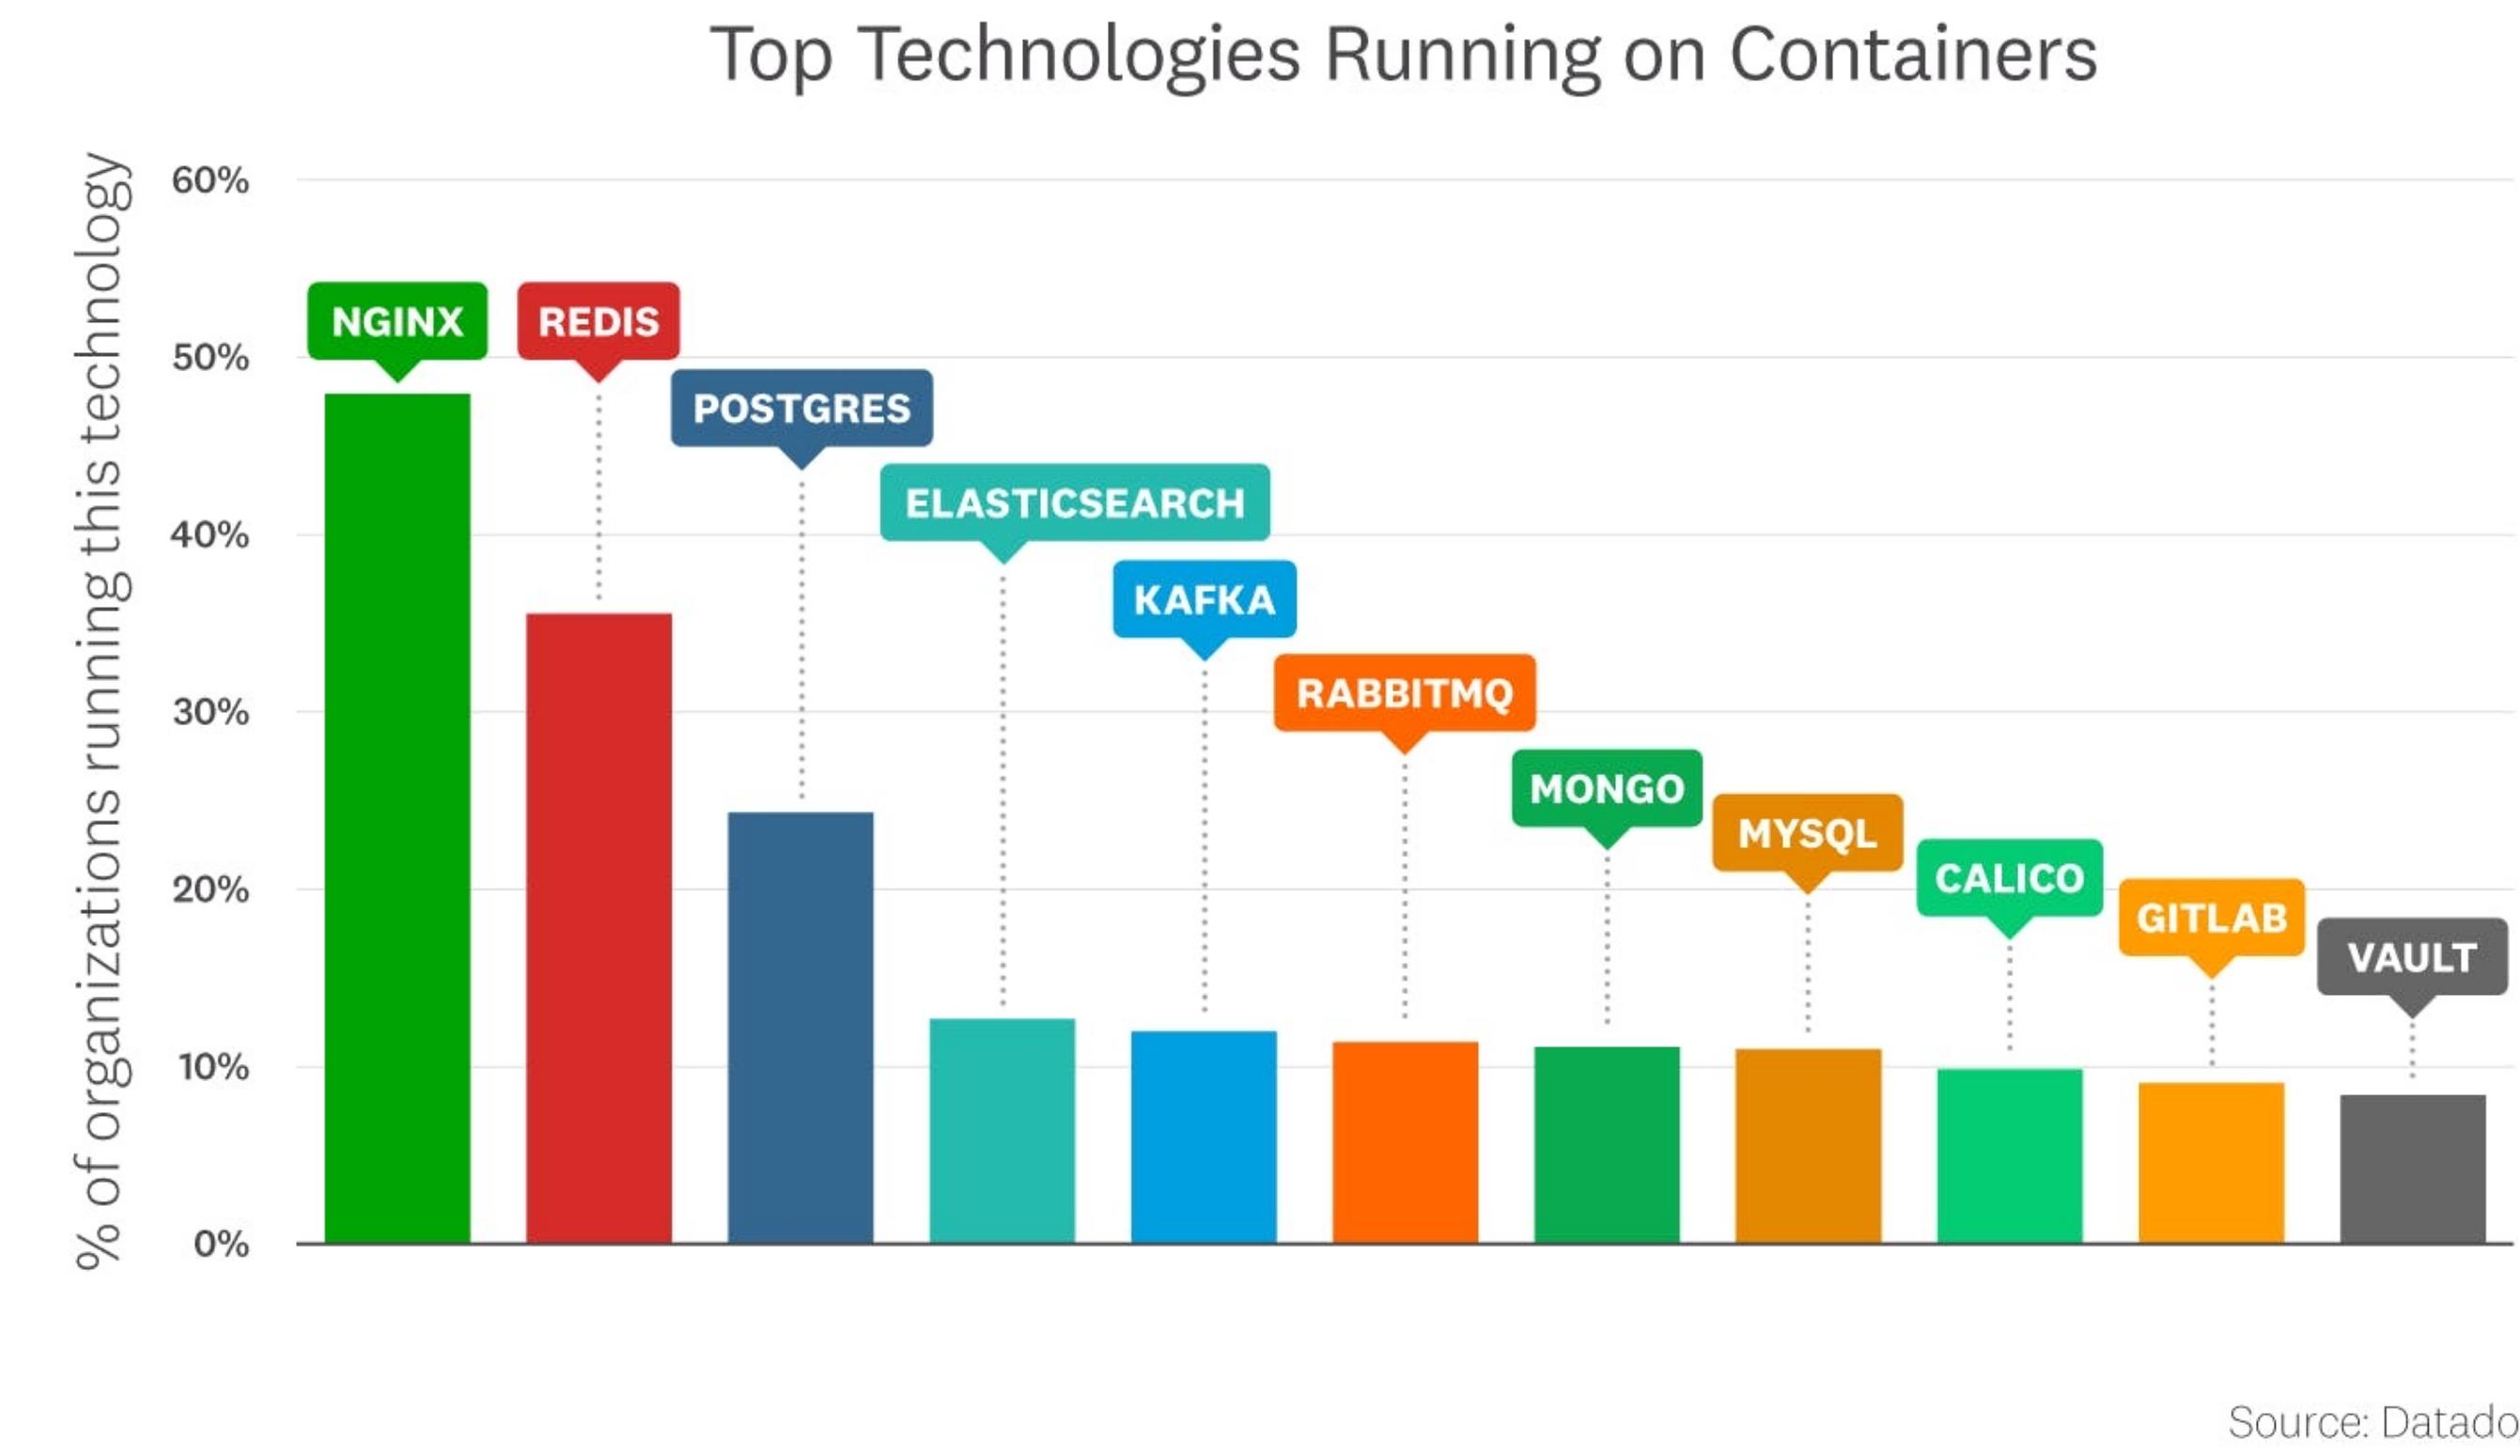


# From Stateless to Stateful



# As Kubernetes has matured

## Stateful Applications are common



[Datadog Real-World Container Use Report – Nov 2022](#)

# Application Consistency Spectrum



## Crash Consistent

- Storage snapshots

## “App” Consistent

- Freeze data service
- Storage snapshot
- Unfreeze data service

## Logical Backup

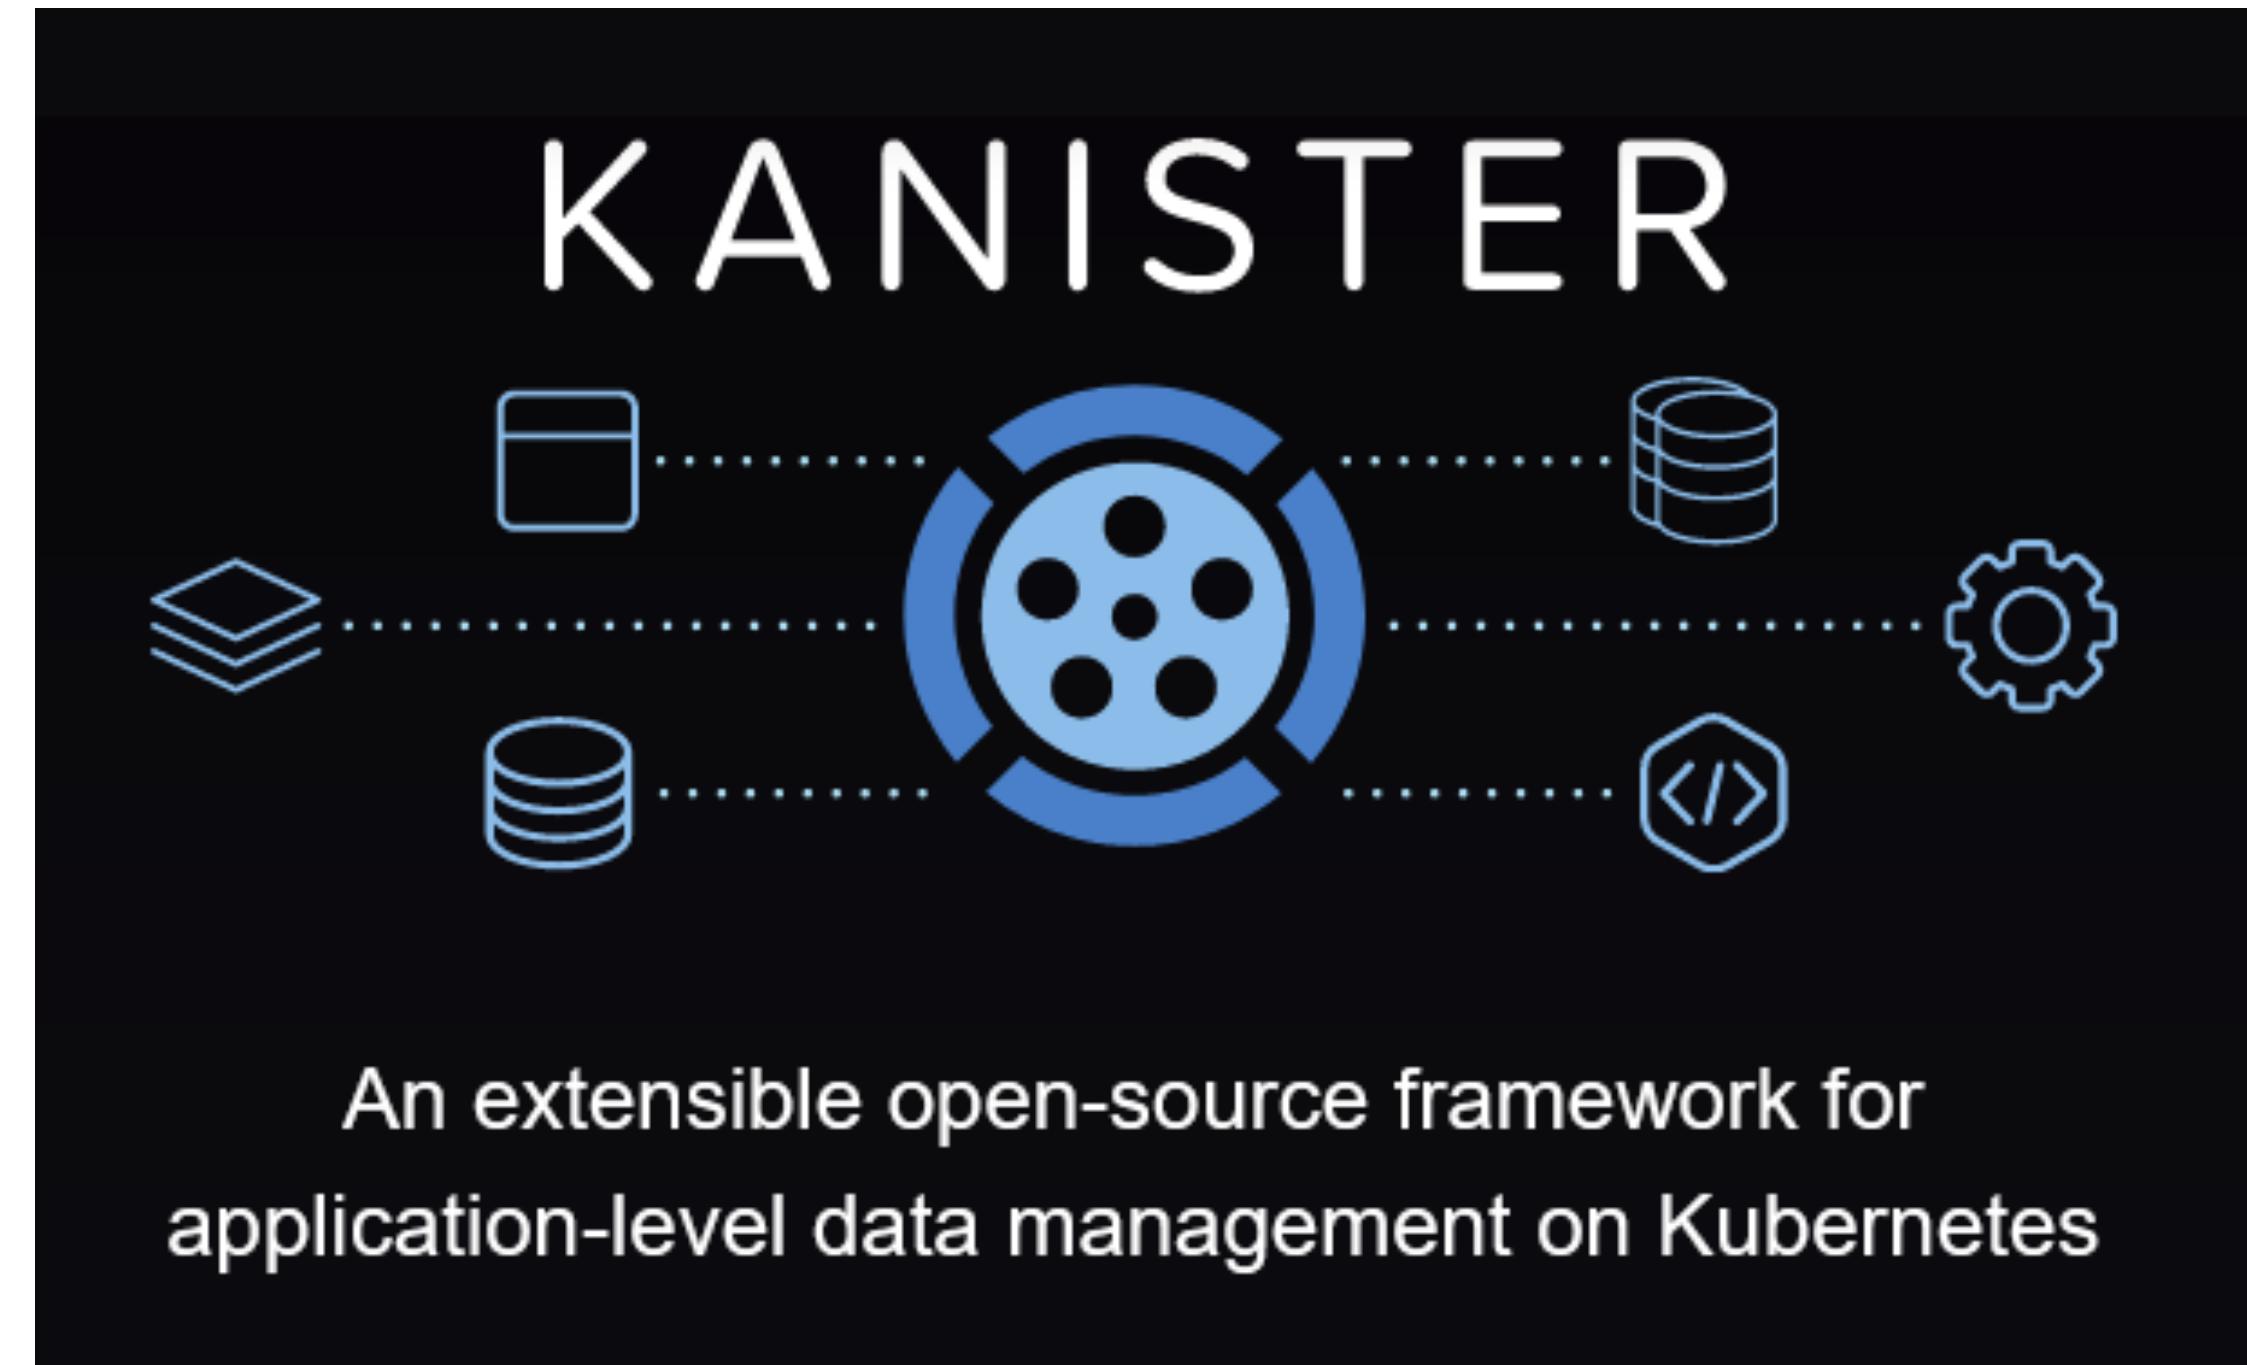
- Logical dumps via data service-specific tool (e.g., `elasticdump`)

## System Consistent

- Full app capture
- Combination of tools across data and storage layers



## Introducing Kanister



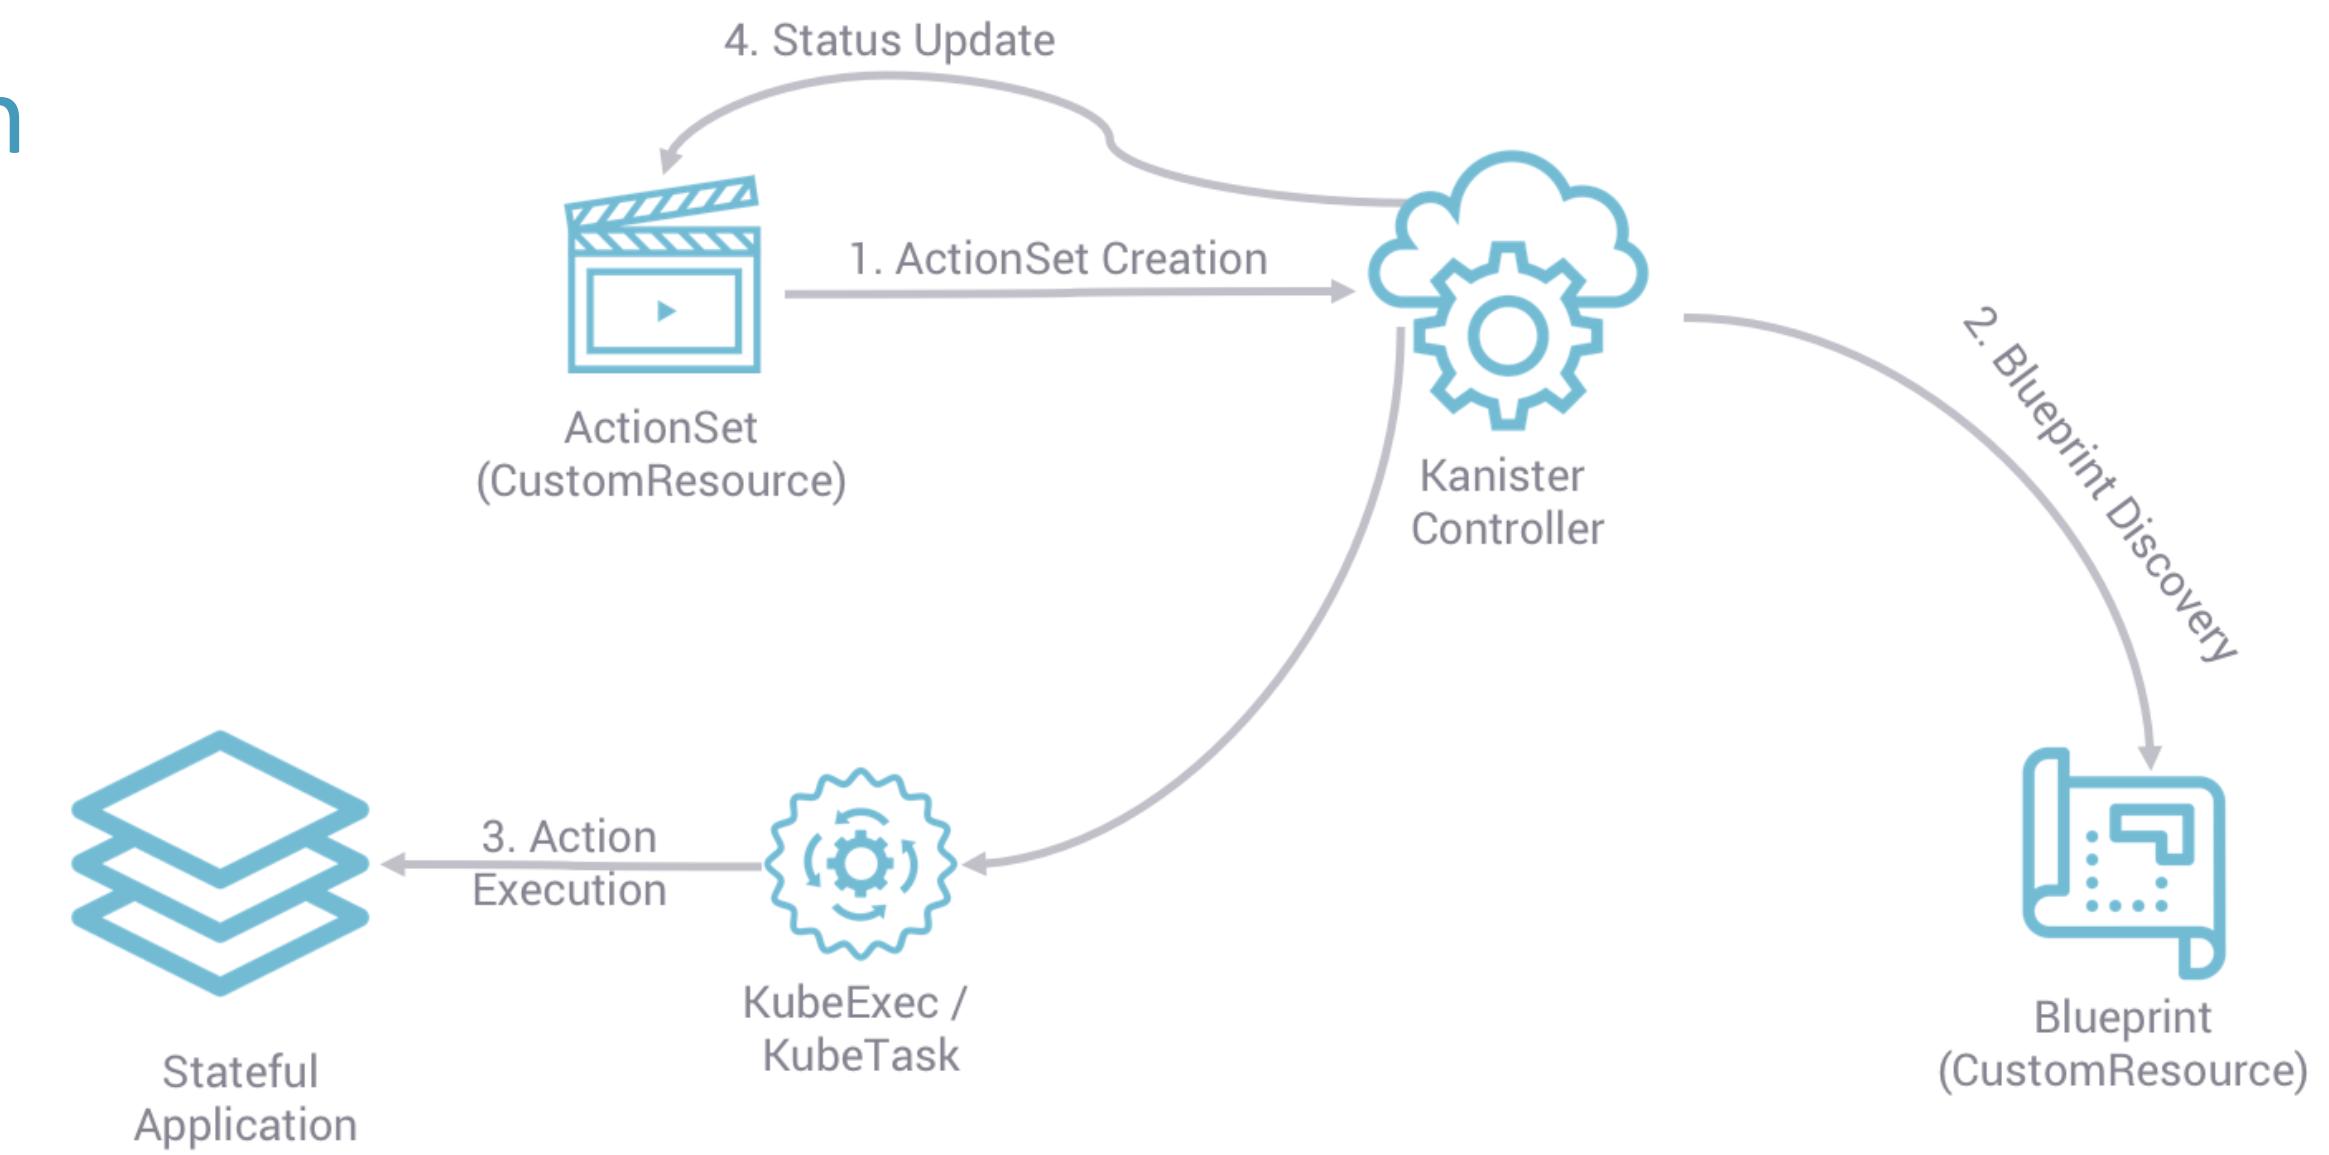
KANISTER allows domain experts to capture application-specific data management tasks in blueprints which can be easily shared and extended.

# How Kanister.io Orchestrates Data Protection

A cloud native, extensible framework for app consistent operations

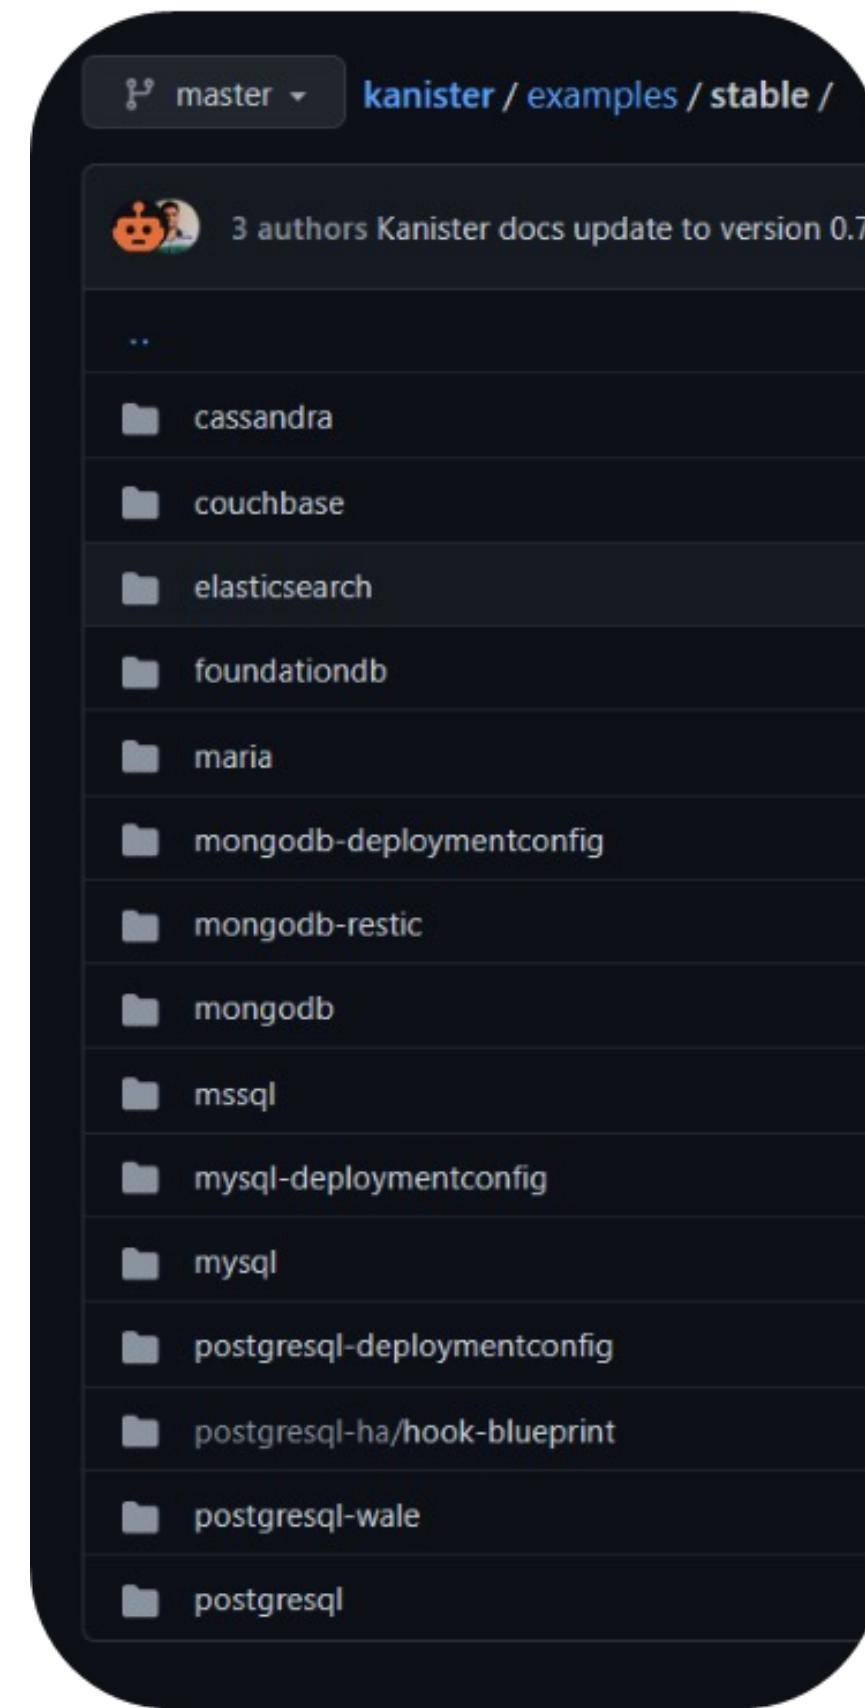
A Kanister controller, installed by Helm chart, provides new Custom Resource Definition

- ActionSet: an instance of a Blueprint operation
- Blueprint: orchestrated actions for backup and recovery
- Profile: infrastructure endpoint config



# Execution Walkthrough

## Kanister Deployment & Blueprints



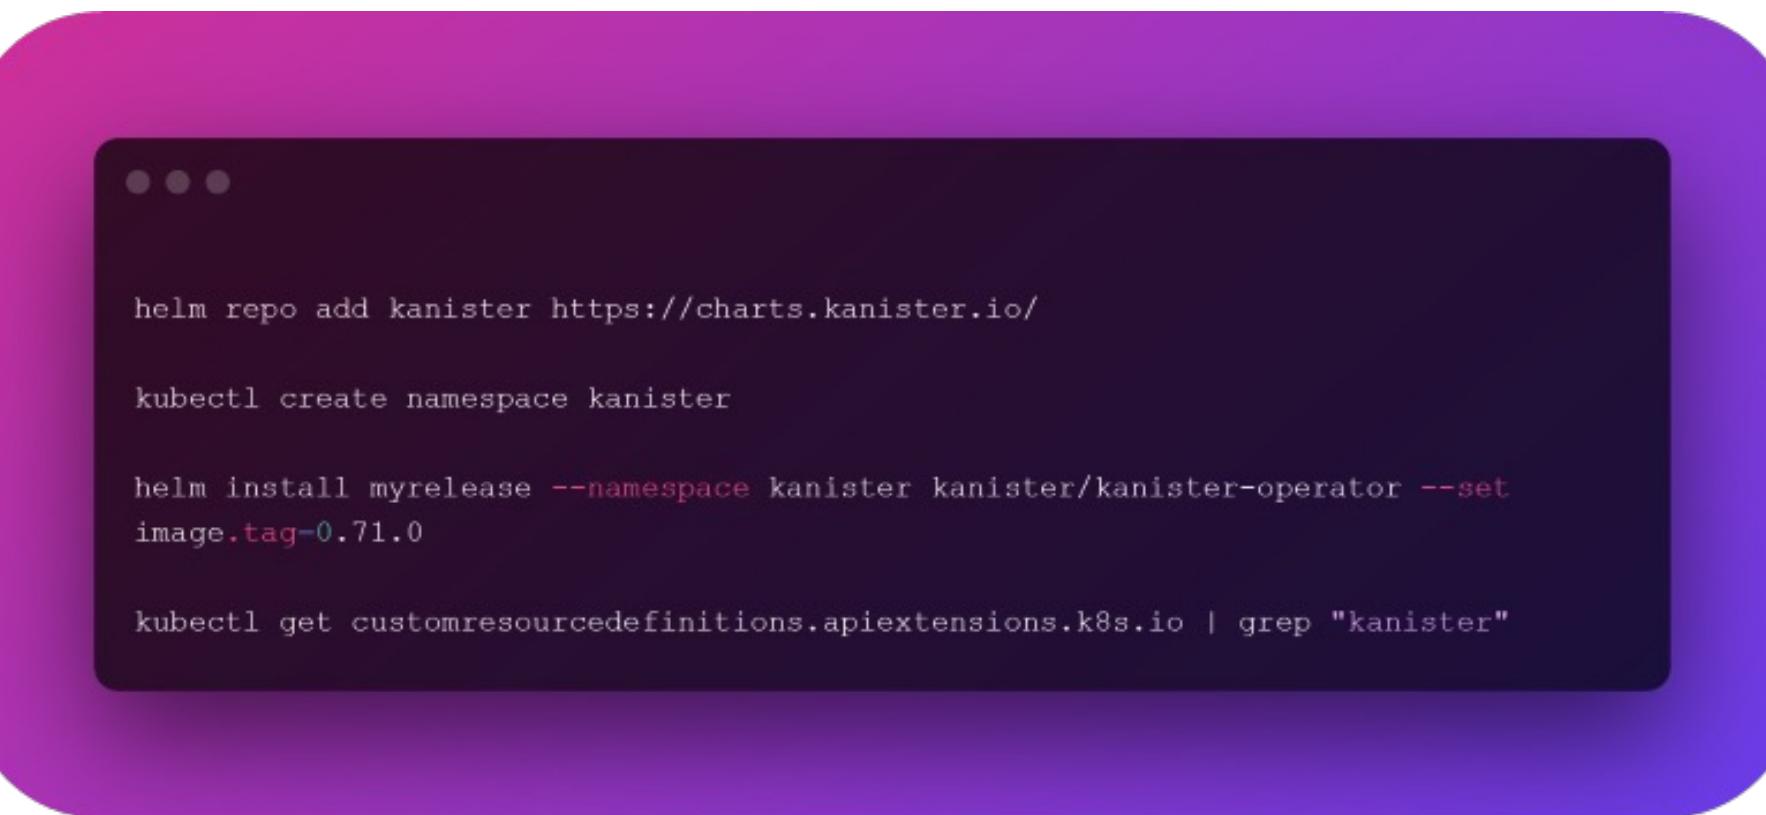
Controller



Blueprint

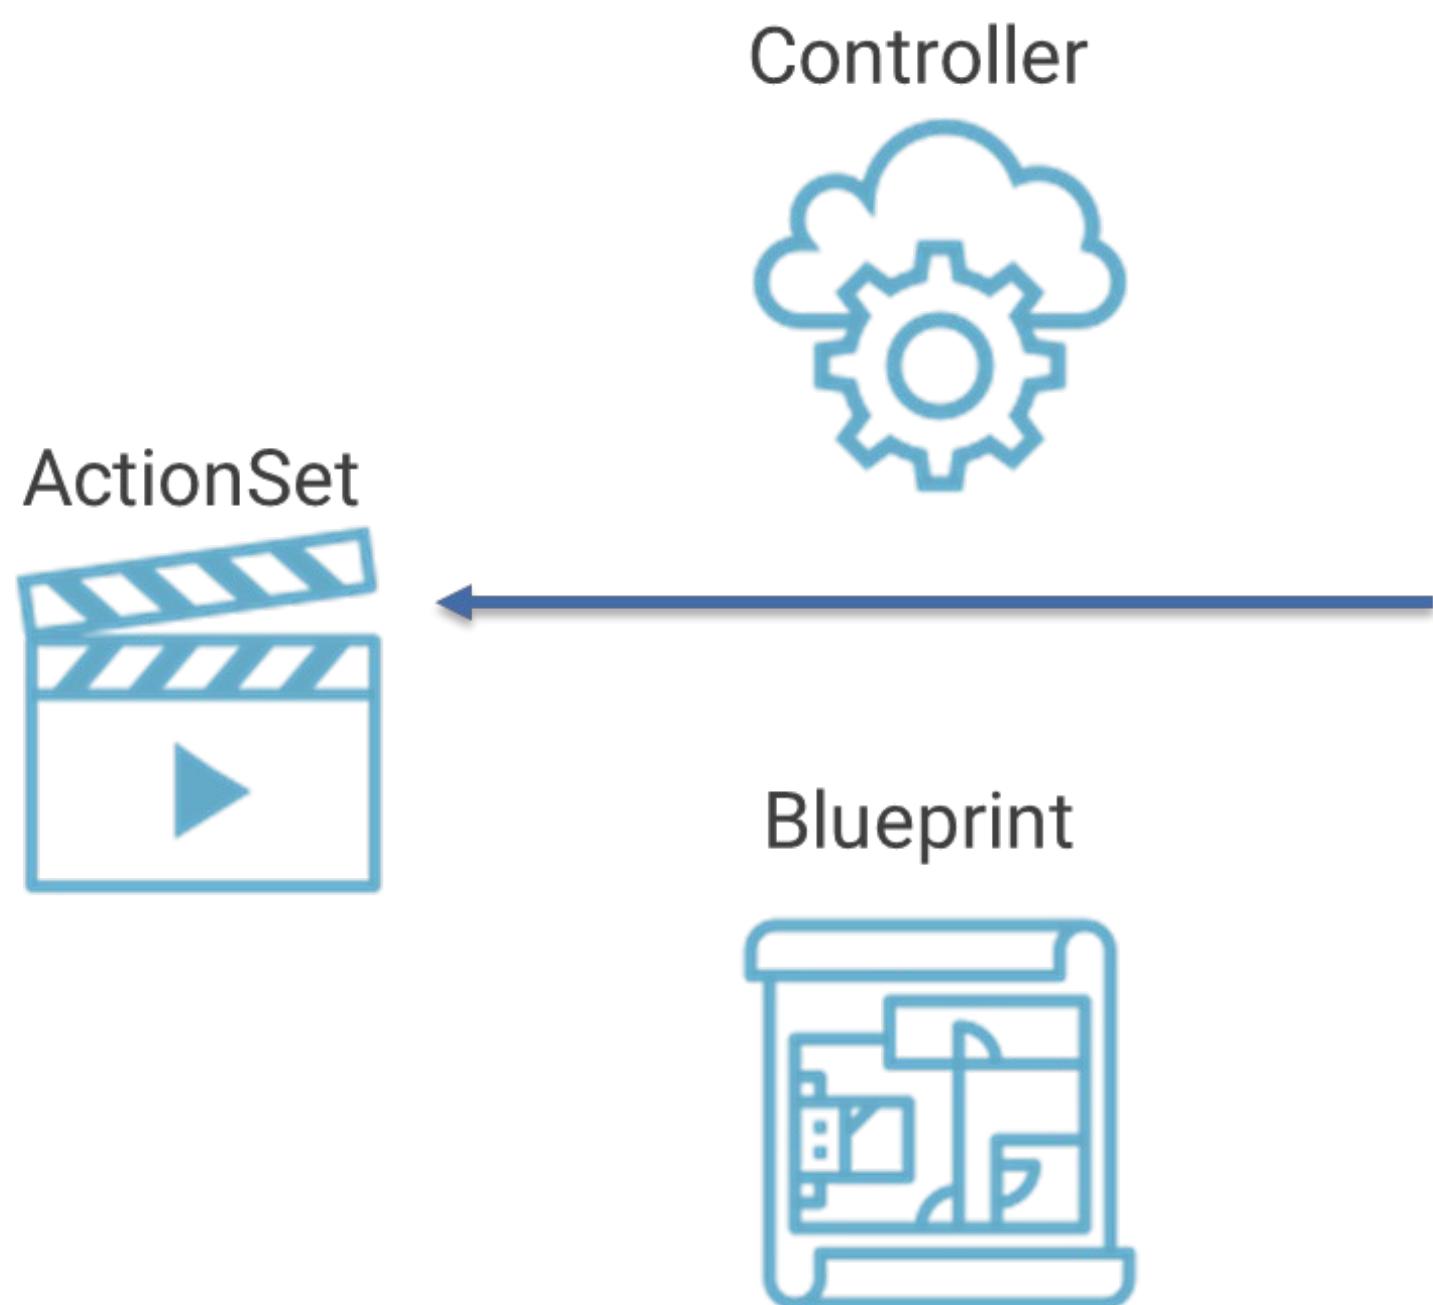


Database Workload



# Execution Walkthrough

## ActionSets



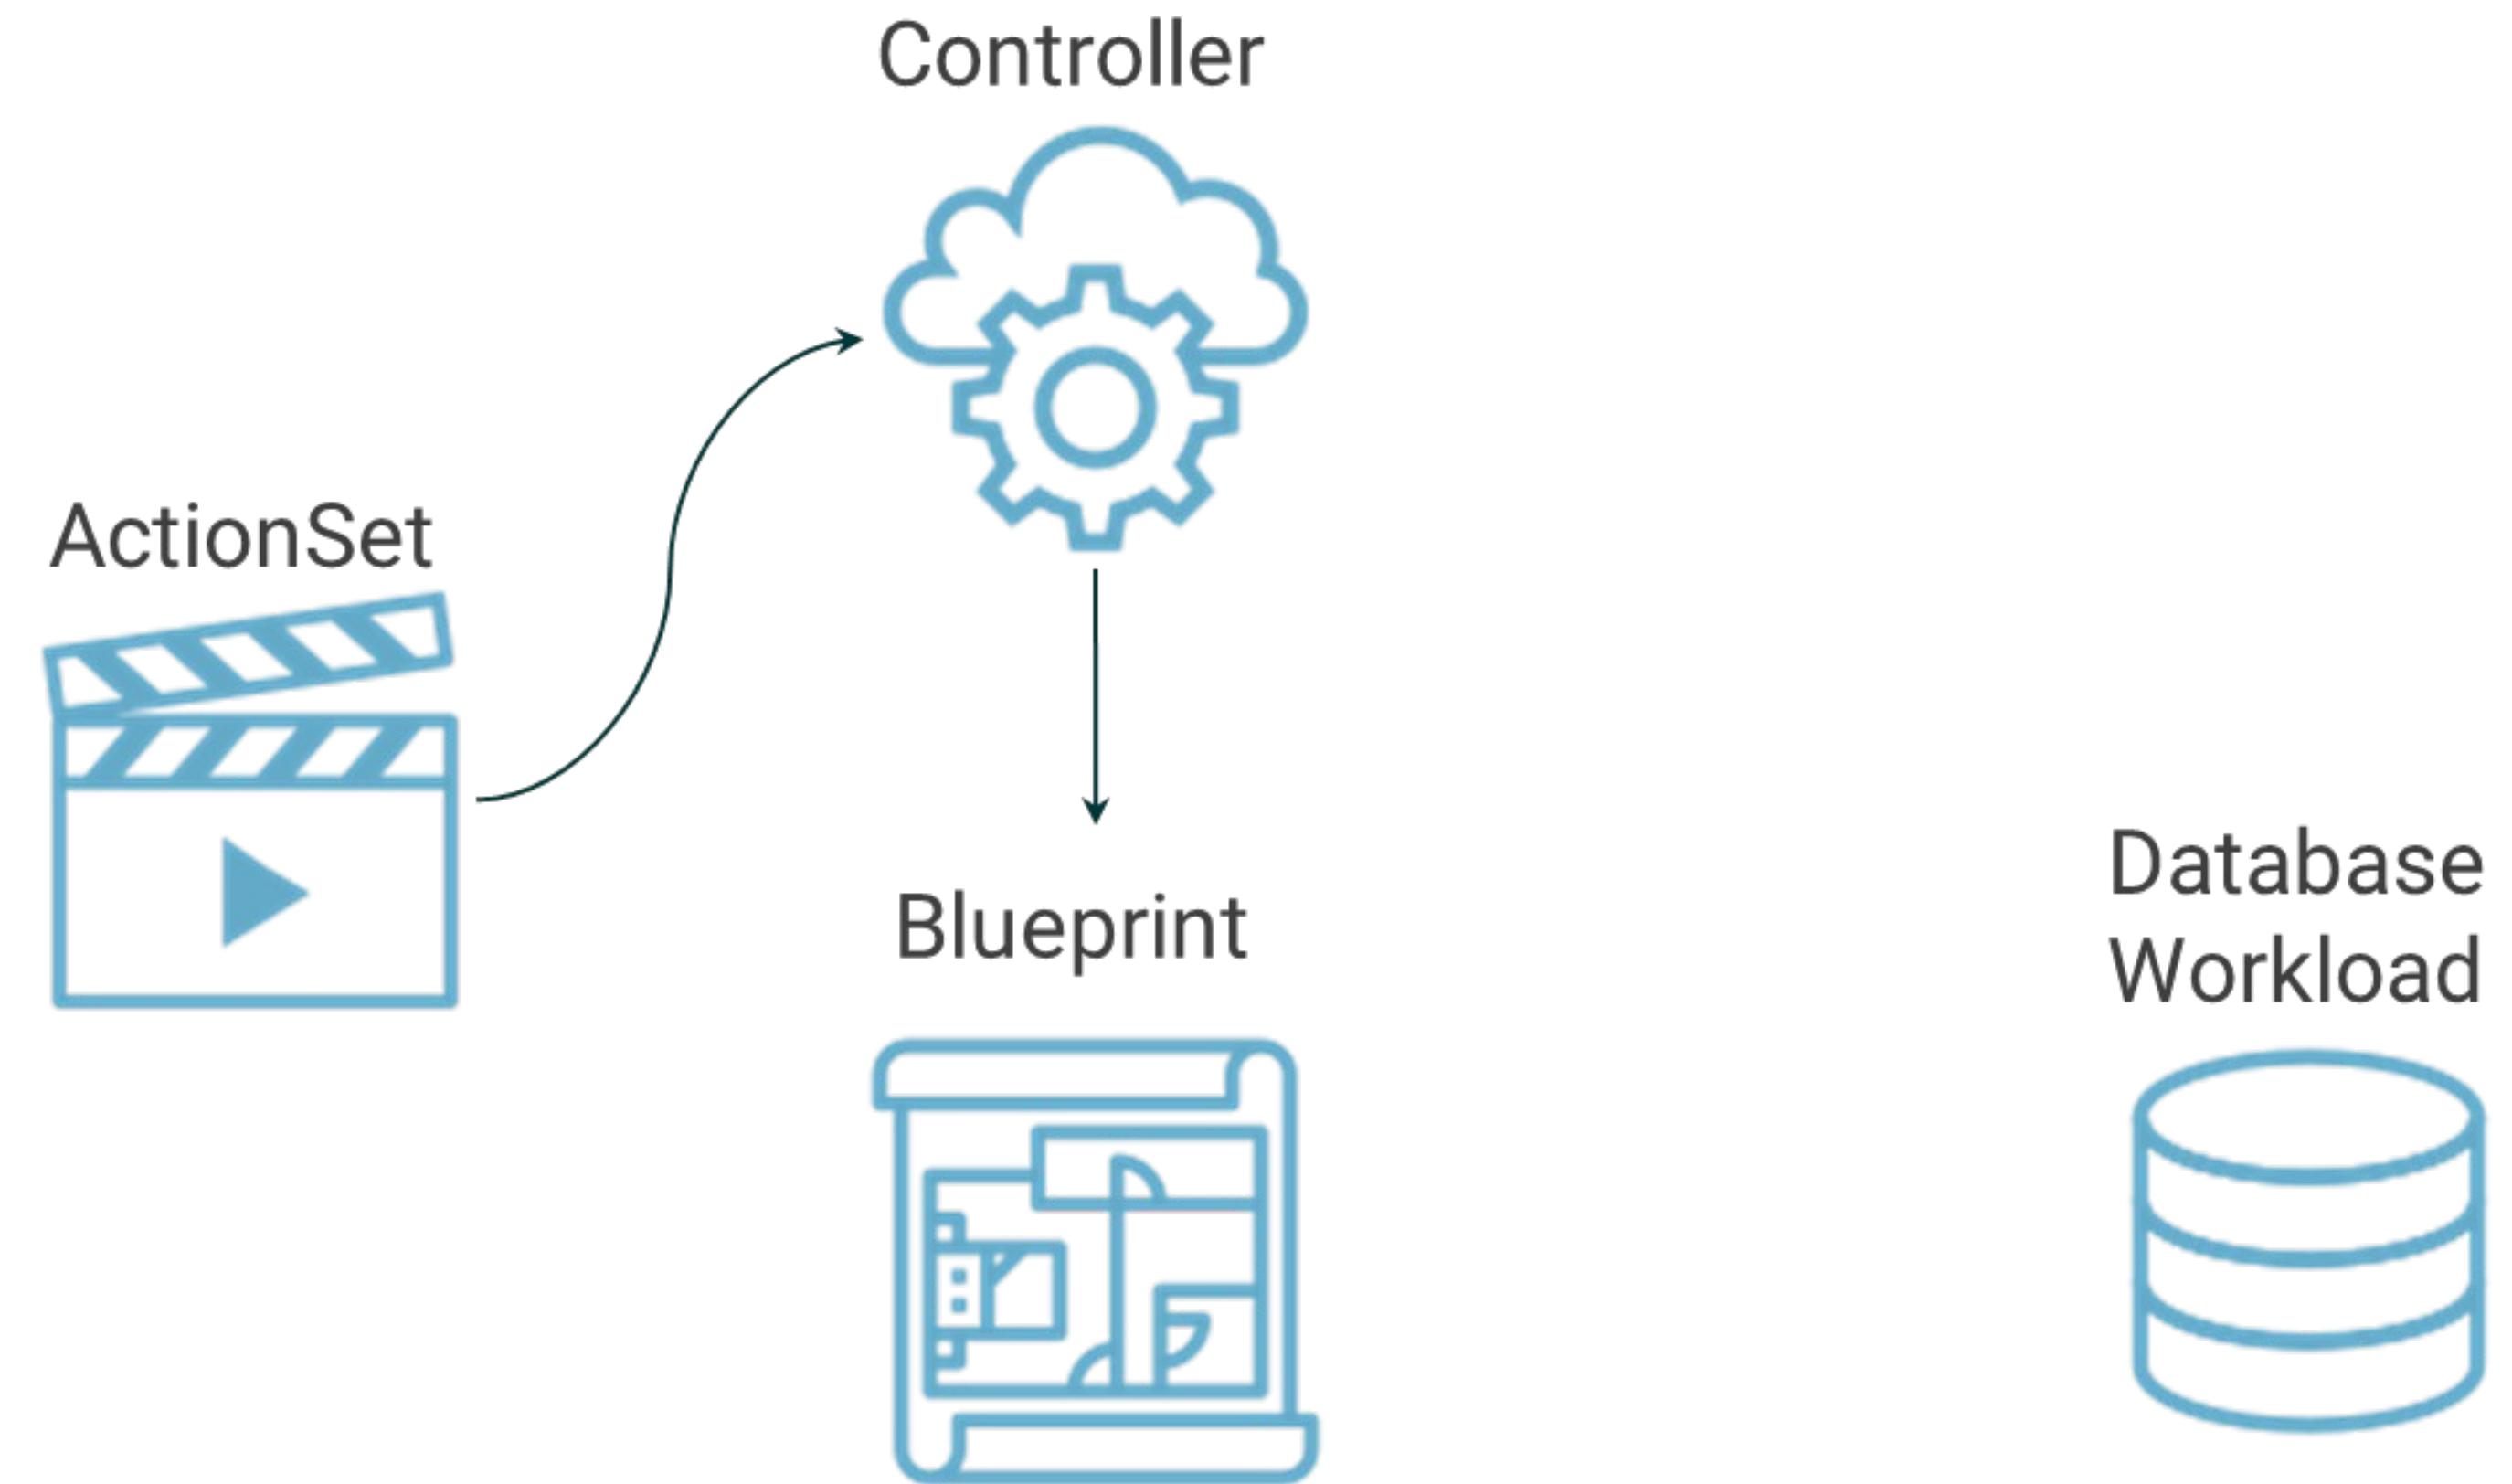
A screenshot of a terminal or code editor displaying a YAML configuration for an ActionSet. The configuration includes metadata like name and namespace, and specifies actions such as a Blueprint for a MySQL backup.

```
● ● ● create an actionset on the project

apiVersion: cr.kanister.io/v1alpha1
kind: ActionSet
metadata:
  name: $backup_name
  namespace: kanister
spec:
  actions:
    - blueprint: mysql-blueprint
      name: backup
      object:
        kind: statefulset
        name: mysql
        namespace: $NAMESPACE
      profile:
        name: $PROFILE
        namespace: kanister
      secrets:
        mysql:
          name: mysql
          namespace: $NAMESPACE
```

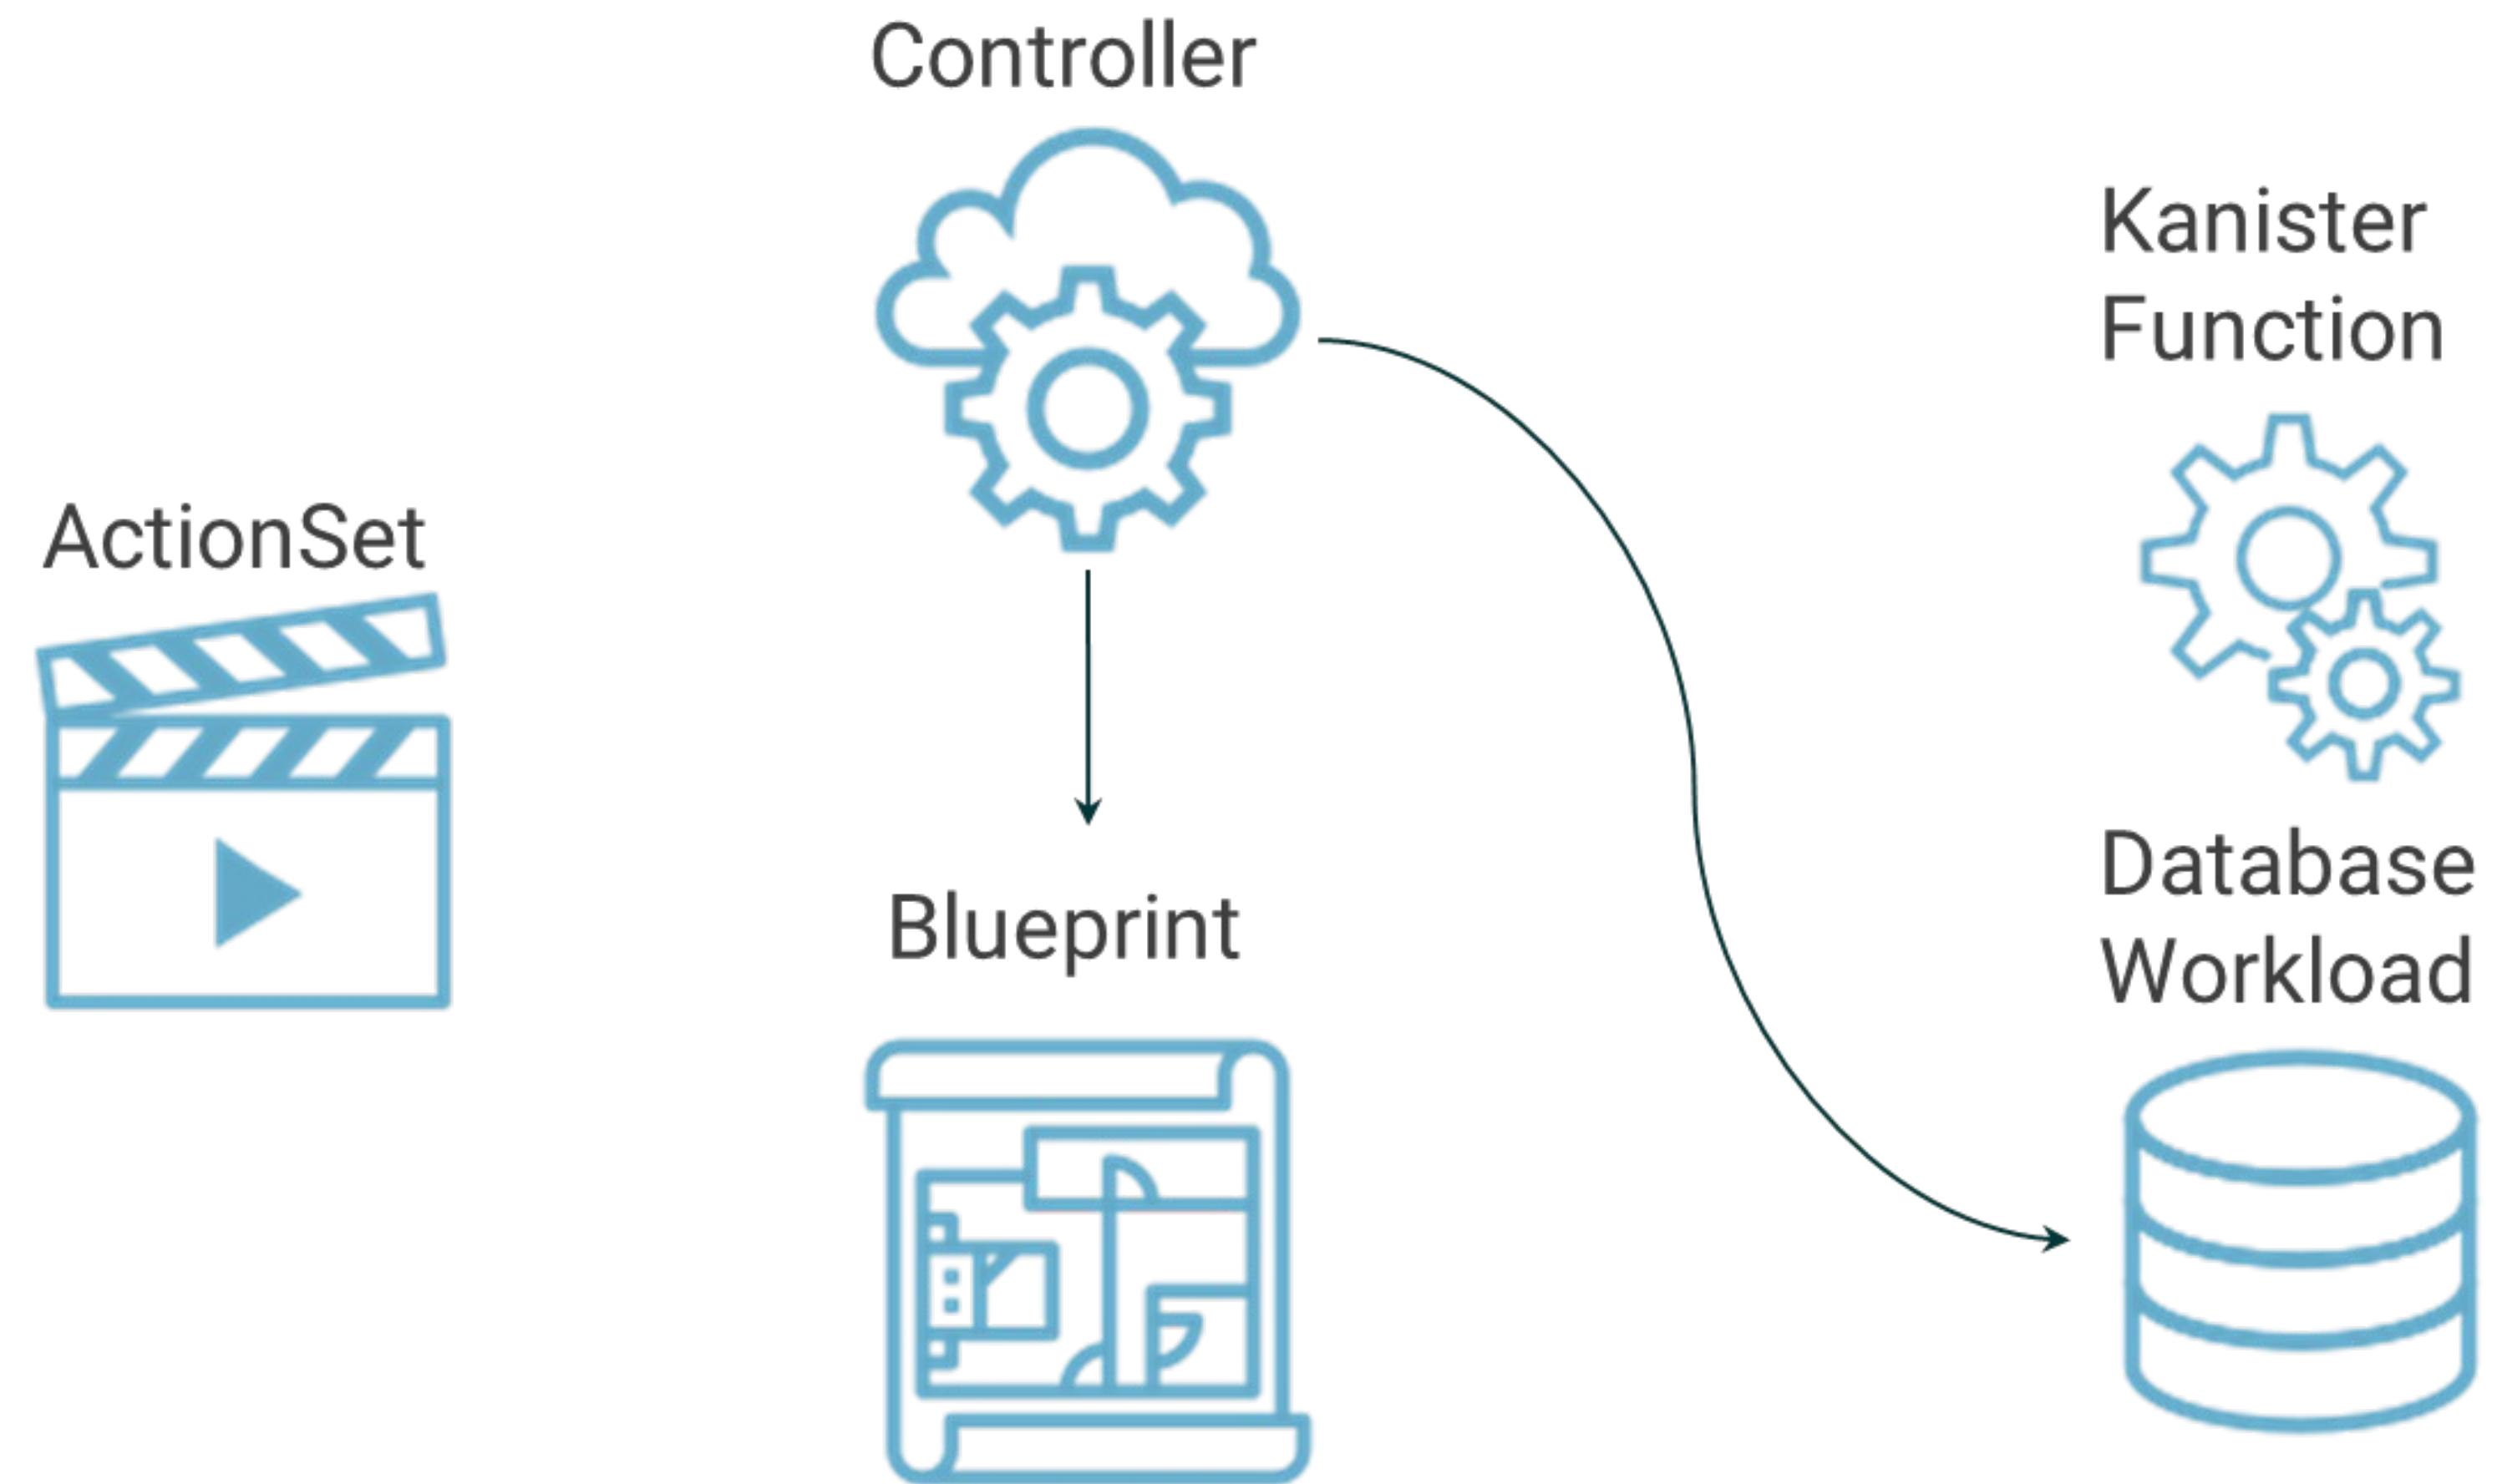
## Execution Walkthrough

ActionSet is started, which then calls the Blueprint



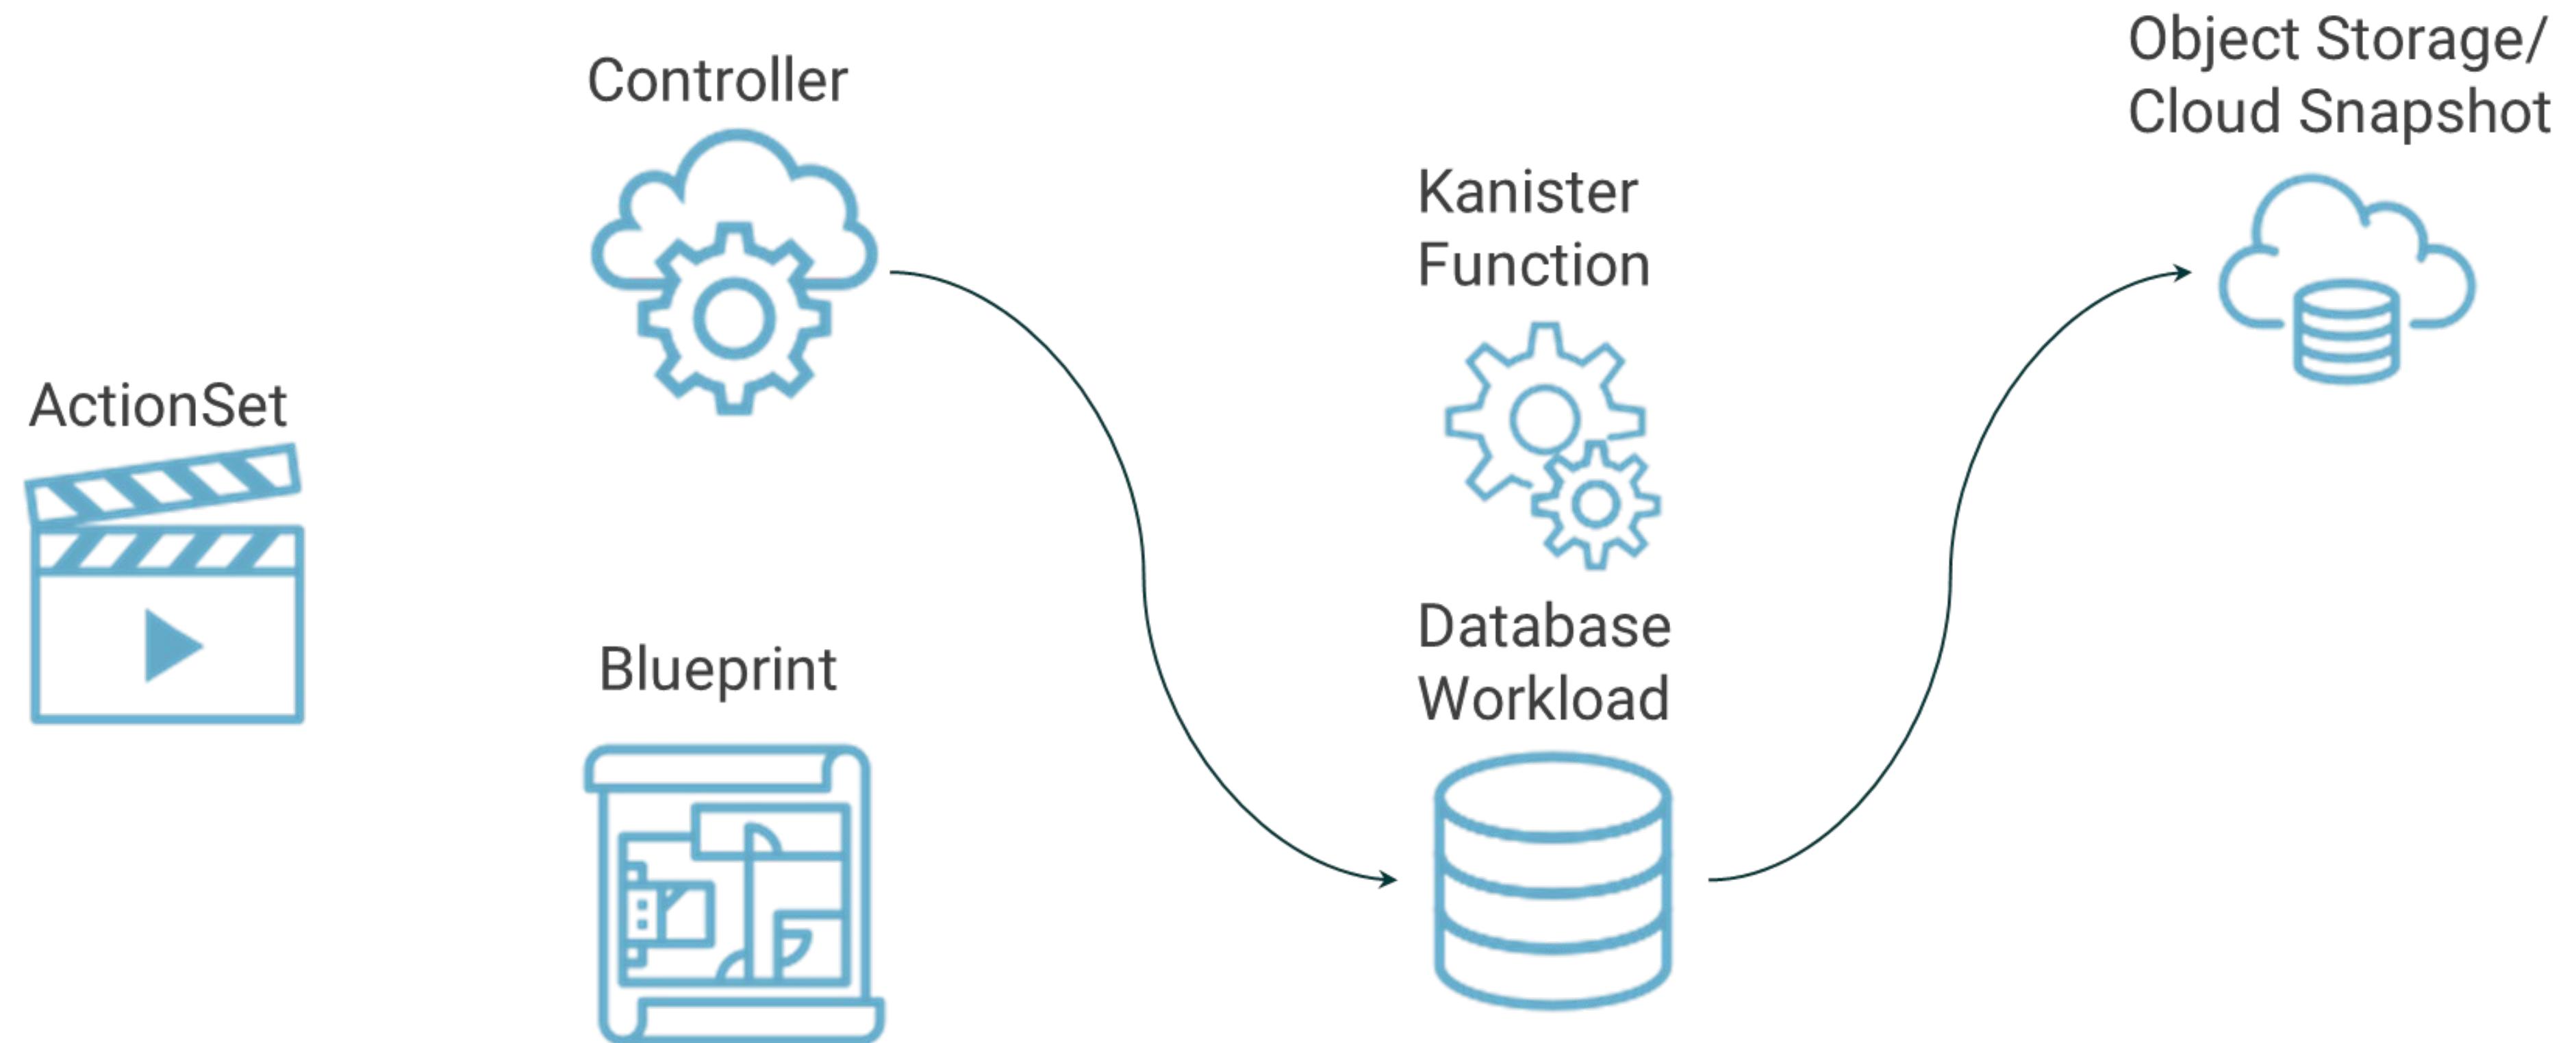
## Execution Walkthrough

Controller then calls Kanister function to interact with Database workload



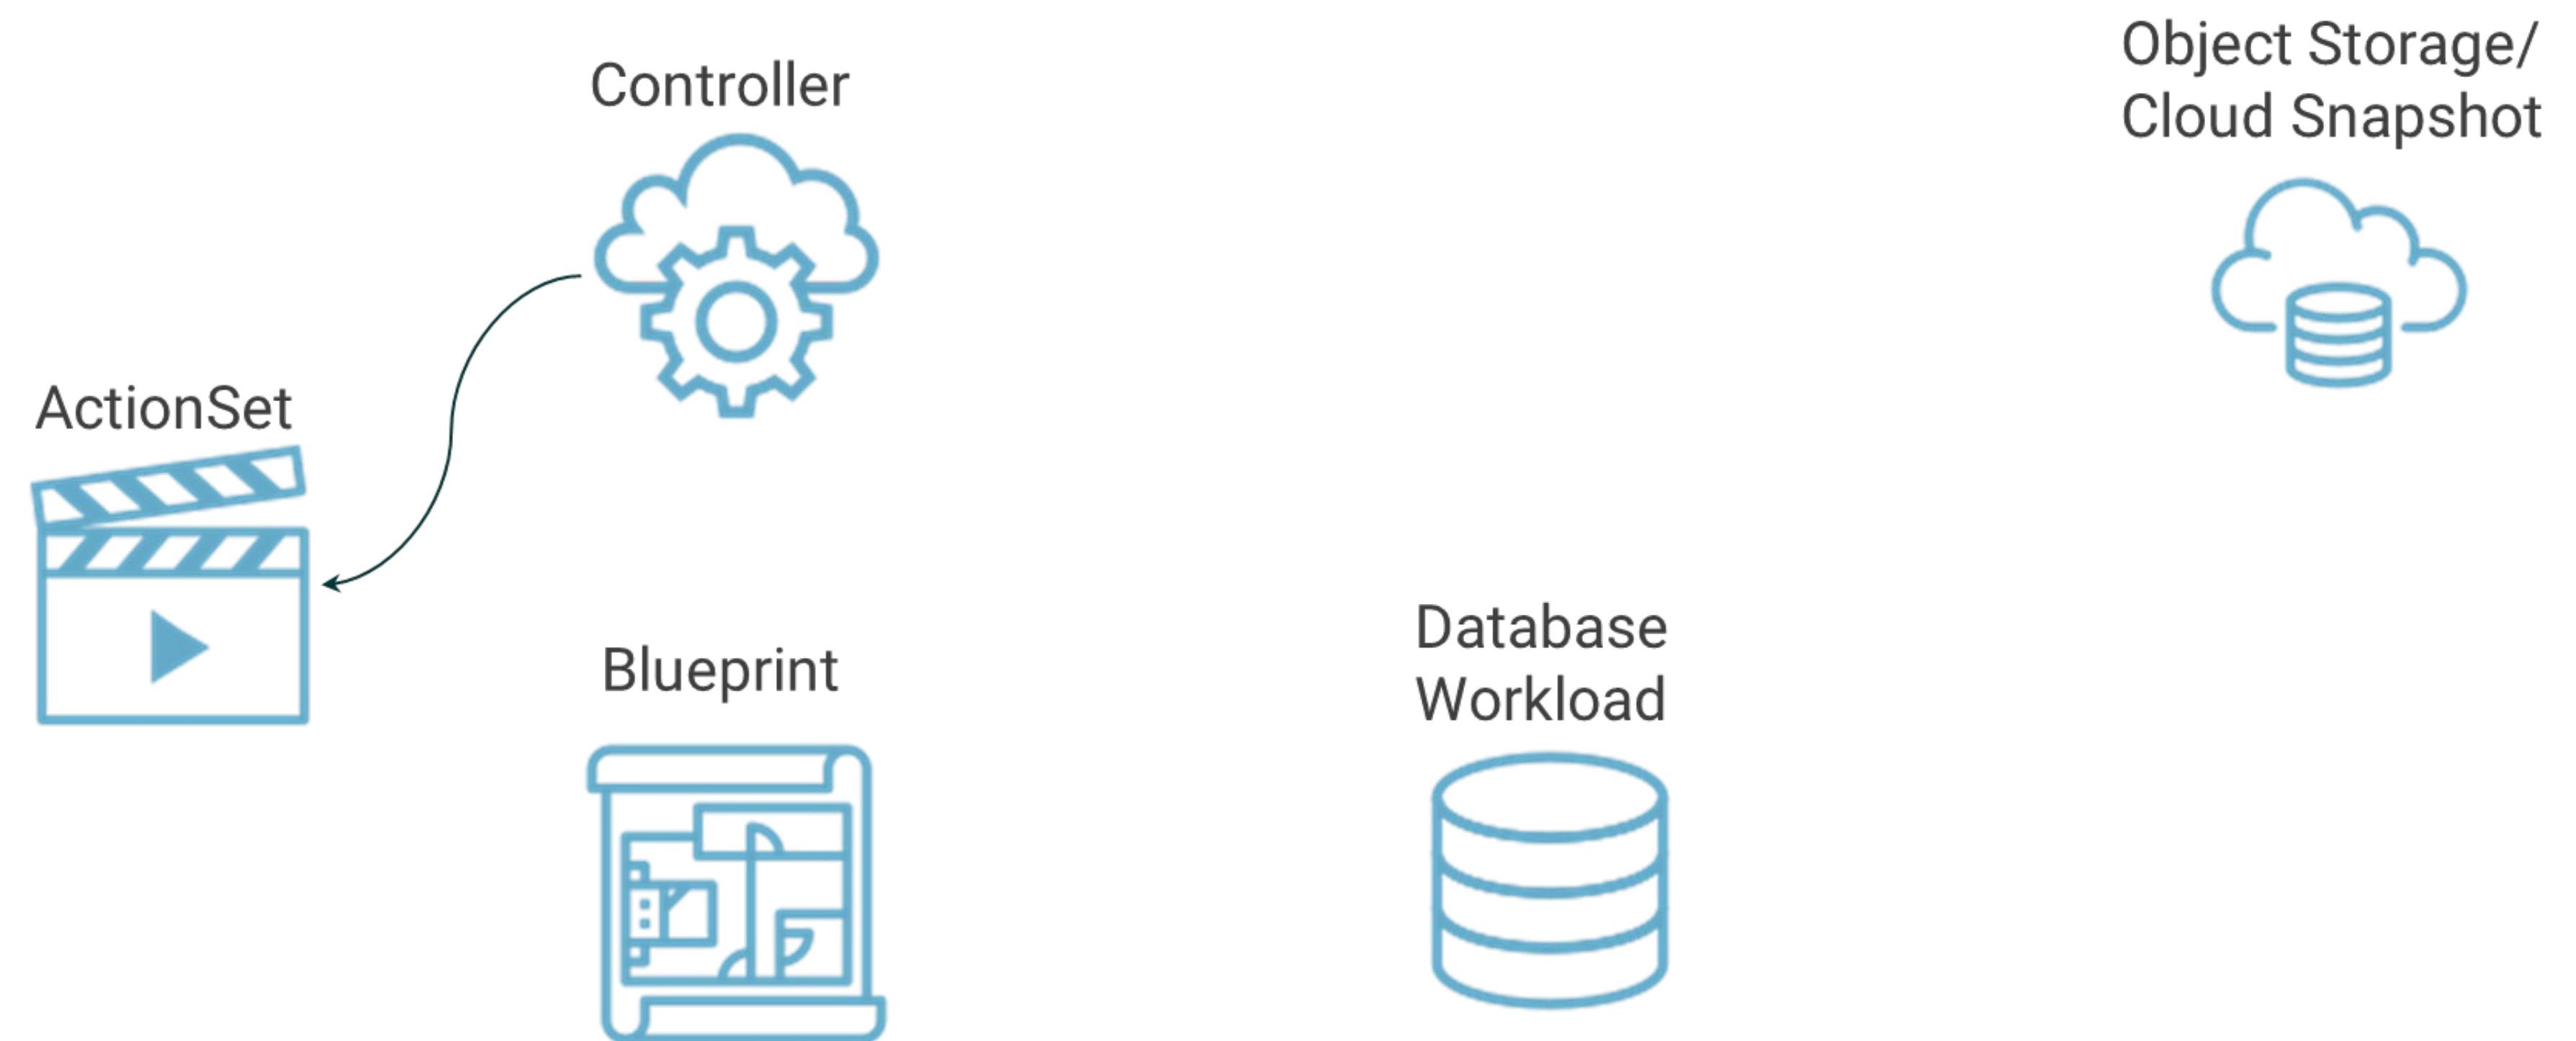
## Execution Walkthrough

When consistency is reached a copy is then sent to backup location



# Execution Walkthrough

Report back status to ActionSet

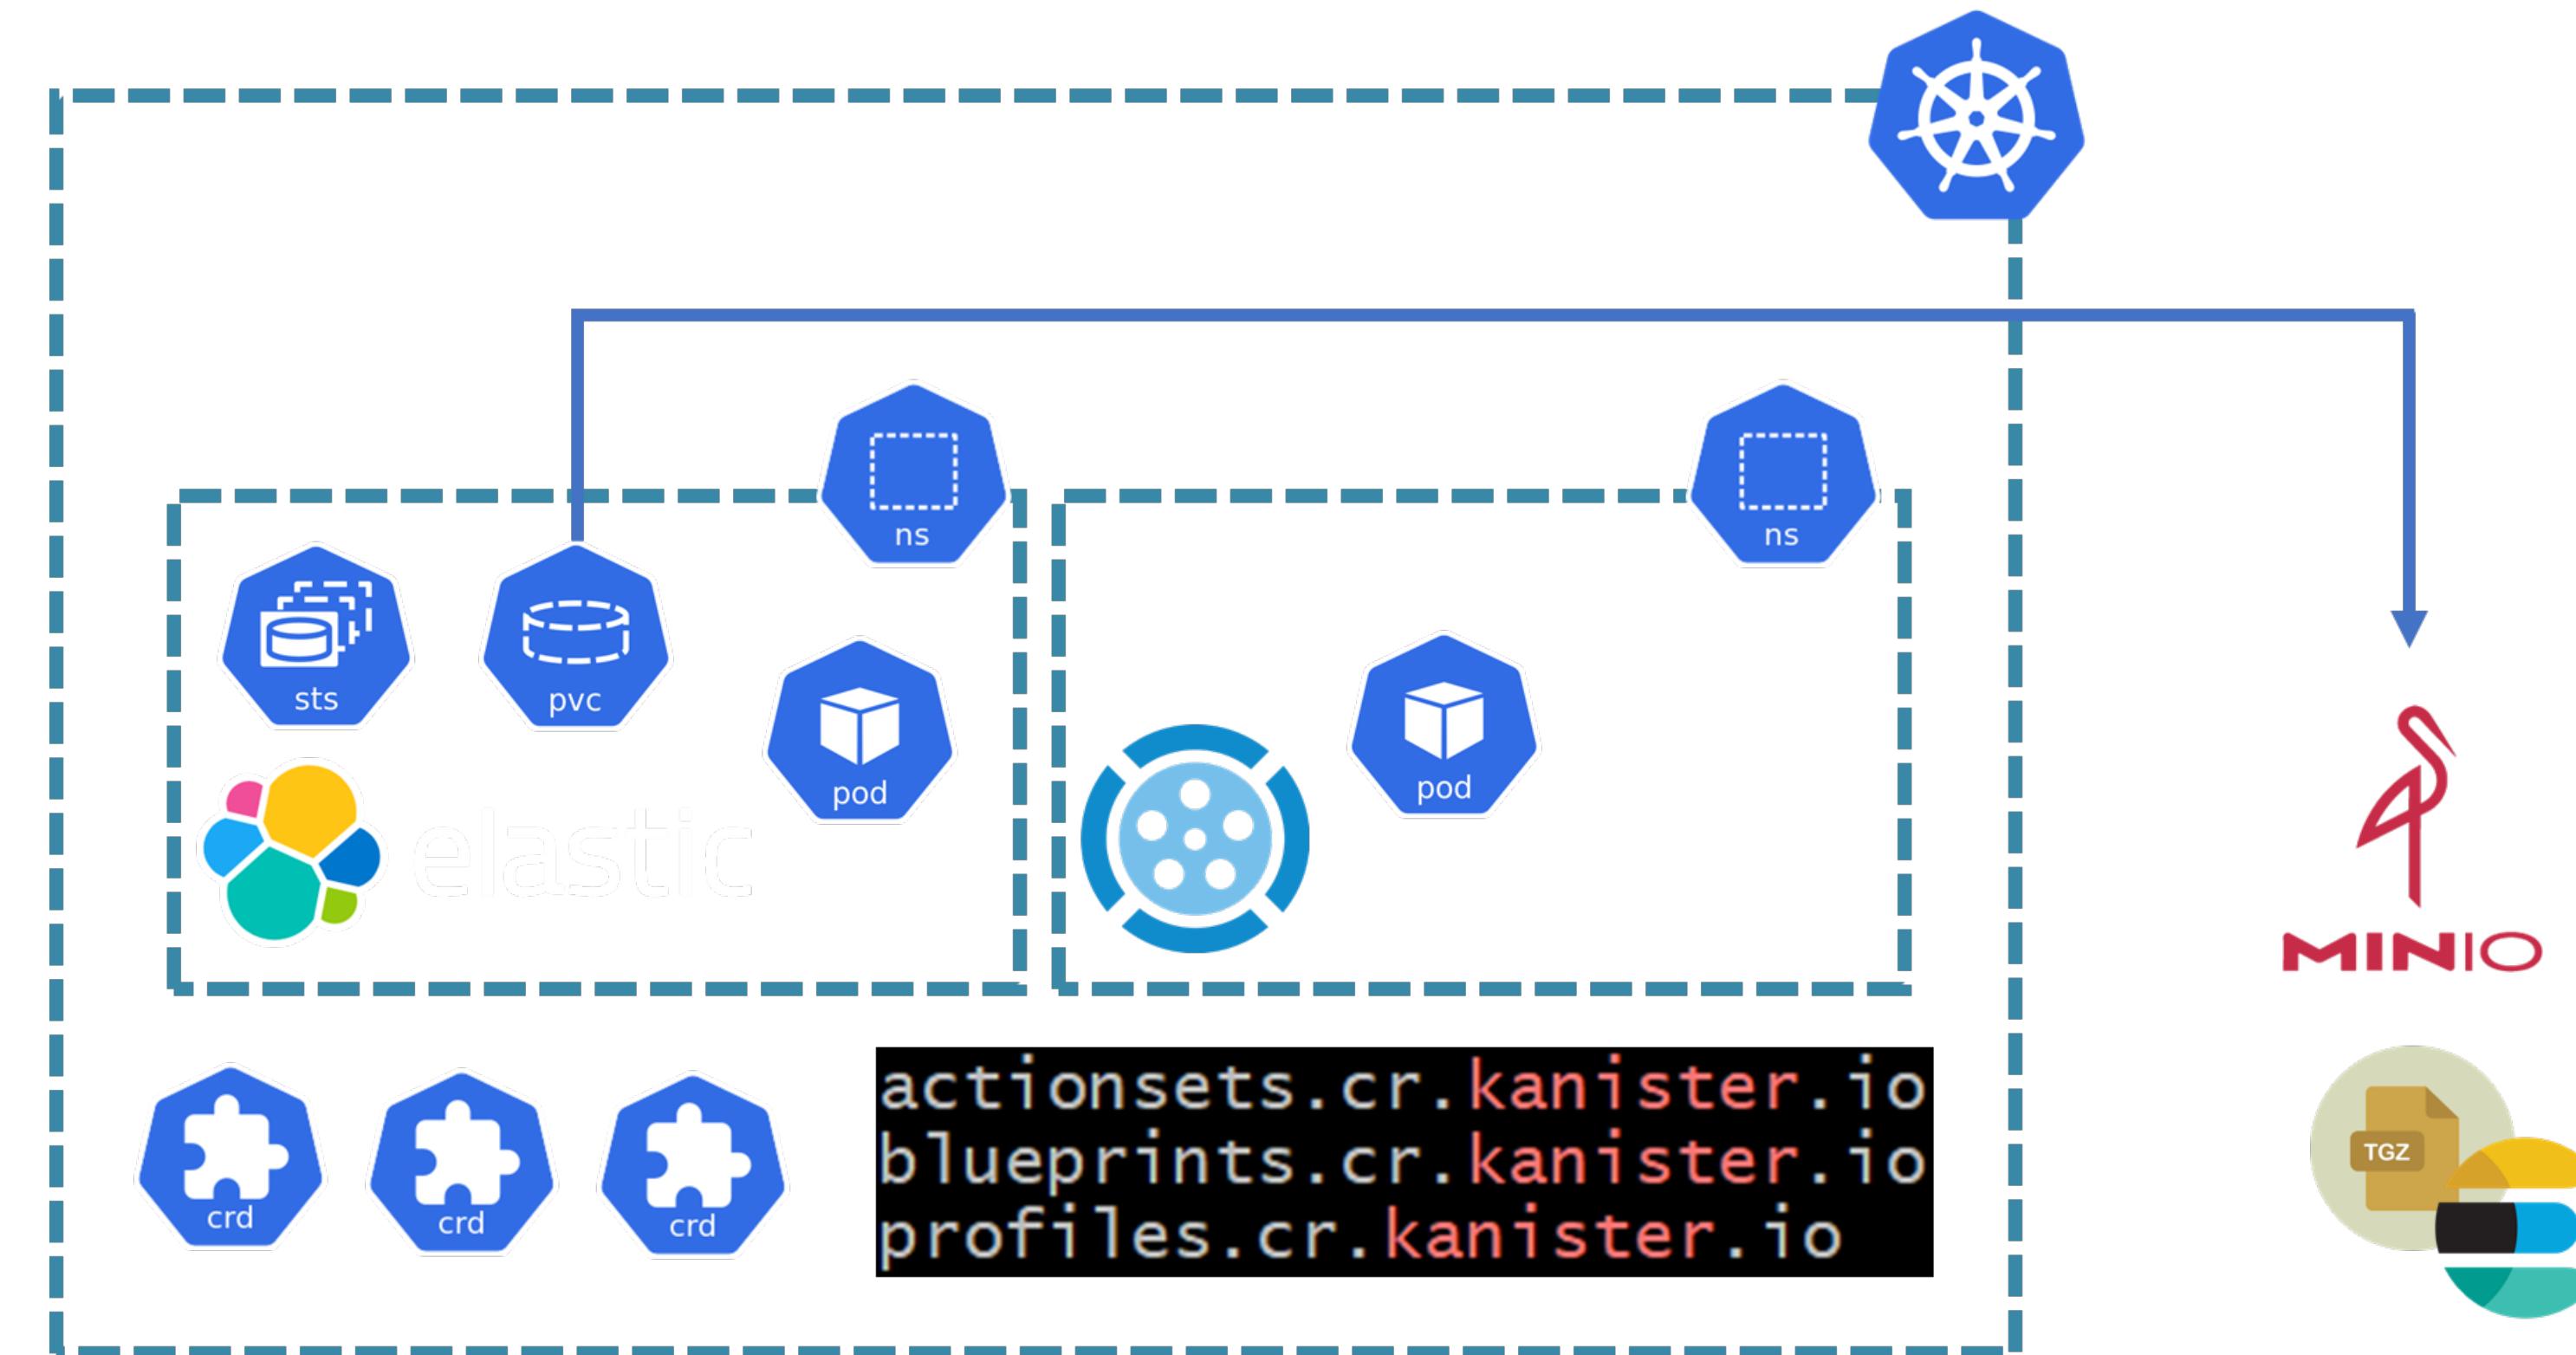


# Demo ElasticSearch with Kanister



# Demo Architecture

Show you what I am about to show you



## **Kahu – What could happen?**

Kahu: Container Data Protection

Kahu is a cloud native tool to backup and restore, perform disaster recovery, and migrate Kubernetes cluster resources and persistent volumes.



Snapshot & Backup of your  
Kubernetes resources  
  
Scheduling & Orchestration



KANISTER

Application-level management  
Blueprint mechanism for  
custom applications

## Resources

Kanister.io is a cloud-native, open source, extensible framework for Kubernetes data protection!

Please adopt and join us at:

Community:  
biweekly Zoom meetings and Slack

Tutorial on <https://KubeCampus.io>

Example blueprints:

<https://github.com/kanisterio/kanister/tree/master/examples>



<https://Kanister.io>

# Getting Started with Kanister

Operators & Developers welcome!

**Kanister is available as an open-source project today!**

 [github.com/kanisterio/kanister](https://github.com/kanisterio/kanister)

 [@kanisterio](https://twitter.com/kanisterio)

 [#kanisterio](#)

 [tiny.cc/kanisterio](https://tiny.cc/kanisterio)

```
...  
  
$ git clone git@github.com:kanisterio/kanister.git  
# install Kanister operator controller  
$ kubectl apply -f bundle.yaml  
# install your application  
$ kubectl apply -f examples/mongo-sidecar/mongo-cluster.yaml  
# use an existing blueprint, tweak one, or create one yourself  
$ kubectl apply -f examples/mongo-sidecar/mongo-blueprint.yaml  
# perform operations (requires setting secrets and configmap)  
$ kubectl create -f examples/mongo-sidecar/backup-actionset.yaml
```

# Enterprise Orchestrated Data Management

Free up to **5 Nodes**  
Fully Featured



[kasten.io/free-kubernetes](https://kasten.io/free-kubernetes)

\*Subscription terms vary

KASTEN  
by Veeam

**THANK YOU!**